

## Setup

1. Login to your Kaggle account.
2. Click on the Code link.
3. Click on the New Notebook button.
4. Change the title of the page on the upper left corner so that it obeys this format: "ML in Python, Spring 2023, Homework 4, Your Name Surname" (e.g. ML in Python, Spring 2023, Homework 4, Zafer Aydın).
5. Click on the "Add Data" button on the upper right corner.
6. Click on the "Search keyword or URL text box" below Add Data. Search for "Housing Prices Competition for Kaggle Learn Users" by entering this text to the search box.
7. Go to page numbered 2.
8. Click on the + button next to the dataset uploaded by user A.P. to add dataset. When you bring your mouse on the + button you should see Add Dataset not Add Notebook Output.
9. Click on the × button next to Add Data at the upper right corner to close the window for adding data.
10. Click on the +Code button at the lower left side of the code cell to add a new code cell.
11. You can start from the template code called ML in Python, Spring 2023, Homework 4, Template (which is made available as ml-in-python-spring-2023-homework-4-template.ipynb).

## Assignment

Implement the steps below on top of the template code. Put the code of each question into a separate cell (e.g. question 1a goes to one cell, question 1b goes to another cell, question 2 goes to another cell, etc.)

1. Define an XGBoost regressor model with default parameters, `n_jobs=-1`, `random_state=0`.
  - (a) What are default values of `n_estimators` and `learning_rate`? You can access default value of `n_estimators` as `print(model.n_estimators)` once you define a model. Include your commands for printing default `n_estimators` into your notebook. The default value of `learning_rate` is given in the tutorial page of XGBoost lesson. Include these default values as markdown cell into your notebook.
  - (b) Train the model on train set (`X_train_imputed`, `y_train`) and compute predictions on test set (`X_test_imputed`). Store the predictions as a pandas data frame, which will have an `Id` column that contains `X_test.index` as the values and a `SalePrice` column that contains predictions on test set as the values. Save this data frame as a csv file such as `submission_default.csv` by setting `index=False`. Submit your predictions to competition. You can find instructions for submitting a csv file to competition in question 4 of homework 2. The link of the competition is <https://www.kaggle.com/competitions/home-data-for-ml-course>. Enter your leaderboard score as a markdown cell to your notebook.
2. Optimize the `n_estimators` and `learning_rate` hyper-parameters of `XGBRegressor` by grid search optimization. Use `GridSearchCV` class of `scikit-learn`, which can be accessed at [https://scikit-learn.org/stable/modules/generated/sklearn.model\\_selection.GridSearchCV.html](https://scikit-learn.org/stable/modules/generated/sklearn.model_selection.GridSearchCV.html). The link includes some code examples for using `GridSearchCV`. Implement the following steps
  - Import `GridSearchCV` from `sklearn.model_selection`.
  - Define a Python list called `n_estimators_values` that contains 100, 500, and 1000 as the values.
  - Define a Python list called `learning_rate_values` that contains 0.01, 0.05, and 0.1 as the values.
  - Define a Python dictionary called `parameters` that contains 'n\_estimators' and 'learning\_rate' as the keys and `n_estimators_values` and `learning_rate_values` as the values of the dictionary.
  - Define a `GridSearchCV` object called `opt` by sending `XGBRegressor` model and the Python dictionary called `parameters` as input to the constructor method called `GridSearchCV`. Note that the default value of 5 is used for `cv` parameter of `GridSearchCV`, which will perform a 5-fold

cross-validation on training set for each hyper-parameter combination in parameter grid once we call the fit method in the next step.

- Call the fit method of GridSearchCV object opt by passing training set as input (X\_train\_imputed and y\_train). This method will optimize the hyper-parameters and once it finishes it will train the XGBRegressor model using the optimum hyper-parameters to make it available for computing predictions.
- Print the optimum n\_estimators and learning\_rate found. You can access the optimums through the best\_params\_ attribute (e.g. opt.best\_params\_). Include these optimums as a markdown cell.
- Call the predict method of GridSearchCV object opt and compute predictions on test set (X\_test\_imputed). Submit your predictions to competition as in question 1b. You can save your csv file as submission\_grid\_search.csv. Enter your leaderboard score as a markdown cell to your notebook.

3. Optimize the n\_estimators and learning\_rate hyper-parameters of XGBRegressor by randomized search optimization. Use RandomizedSearchCV class of scikit-learn, which can be accessed at [https://scikit-learn.org/stable/modules/generated/sklearn.model\\_selection.RandomizedSearchCV.html#sklearn.model\\_selection.RandomizedSearchCV](https://scikit-learn.org/stable/modules/generated/sklearn.model_selection.RandomizedSearchCV.html#sklearn.model_selection.RandomizedSearchCV).

The link includes some code examples for using RandomizedSearchCV. Implement the following steps

- Import RandomizedSearchCV from sklearn.model\_selection, randint from scipy.stats, and loguniform from scipy.stats.
- Define a Python dictionary called distributions as in the code examples for randomized search. Set n\_estimators to randint(100,1001) and learning\_rate to loguniform(0.01, 0.1). This way randomized search algorithm will sample n\_estimators as a random integer in range 100 to 1000 from randint distribution and learning\_rate as a real-value in range 0.01 to 0.1 from loguniform distribution.
- Define a RandomizedSearchCV object called opt by sending XGBRegressor model and the Python dictionary called distributions as input to the constructor method called RandomizedSearchCV and by setting n\_iter to 9, n\_jobs to -1, and random\_state to 0. This will make a total of 9 iterations for randomized search (a total of 9 hyper-parameter combinations will be tried) and will perform a 5-fold cross-validation on training set in each iteration.
- Call the fit method of RandomizedSearchCV object opt by passing training set as input (X\_train\_imputed and y\_train). This method will optimize the hyper-parameters and once it finishes it will train the XGBRegressor model using the optimum hyper-parameters to make it available for computing predictions.
- Print the optimum n\_estimators and learning\_rate found. You can access the optimums through the best\_params\_ attribute (e.g. opt.best\_params\_). Include these optimums as a markdown cell.
- Call the predict method of RandomizedSearchCV object opt and compute predictions on test set (X\_test\_imputed). Submit your predictions to competition as in question 1b. You can save your csv file as submission\_randomized\_search.csv. Enter your leaderboard score as a markdown cell to your notebook.

4. Define an XGBRegressor model with n\_estimators set to 1000, learning\_rate to 0.05, n\_jobs to -1, and random\_state to 0. Train the model on the training set (X\_train\_imputed, y\_train) and compute predictions on test set (X\_test\_imputed). Convert the predictions to a pandas data frame and save it as a csv file called submission\_selected.csv. Submit your csv file to competition as in question 1b. Enter your leaderboard score as a markdown cell to your notebook.

5. Import cross\_val\_score from sklearn.model\_selection. Implement a Python function called score\_dataset that performs the following

- Receives train set (i.e. X\_train\_set and y\_train\_set), n\_estimators, and learning\_rate parameters of XGBoost as input.

- Sets the `n_estimators` parameter of XGBoost to `n_estimators` that is input to your Python function, `learning_rate` parameter of XGBoost to `learning_rate` that is input to your Python function, `n_jobs` to -1 and `random_state` to 0.
- Performs a 5-fold cross-validation on training set (`X_train_set` and `y_train_set`) using `cross_val_score` and XGBoost regressor as the model.
- Computes and returns the average mean absolute error (MAE) as the output obtained as the average of five scores from each fold of cross-validation.

6. Compute and print the average MAE scores by calling the `score_dataset` function (i.e. by performing 5-fold cross-validation) using the training set (`X_train_imputed` and `y_train`) for the following scenarios

- `n_estimators` and `learning_rate` set to defaults as in question 1
- `n_estimators` and `learning_rate` set to optimums found by grid search as in question 2
- `n_estimators` and `learning_rate` set to optimums found by randomized search as in question 3
- `n_estimators` and `learning_rate` set to values used in question 4

Include these MAE scores as a markdown cell to your notebook.

7. Fill the table below that includes your cross-validation (computed on training set) and test set scores and include to your solution document (e.g. a Word file).

<code>n_estimators</code> and <code>learning_rate</code>	CV MAE on Training Set	Test Set MAE
Defaults		
Optimized by grid search		
Optimized by randomized search		
<code>n_estimators=1000</code> , <code>learning_rate=0.05</code>		

8. Include your answers for the following questions to your solution document (e.g. a Word file).

- (a) Which approach gives the best cross-validation MAE score? Which approach gives the best test set score (i.e. leaderboard score)?
- (b) Can we say that optimizing hyper-parameters gives better results than using default values?
- (c) If the approach that gives the best cross-validation MAE score is different from the approach that gives the best leaderboard score what can be the reason for this? Hint: Can you consider the training set as small?

## Submission

Once you finish, click File and Download notebook. Submit your notebook with `.ipynb` extension to Canvas. Submit your solutions to questions 7 and 8 as a Word file to Canvas.